Image Caption Generation using CNN-Based Visual Features and Transformer Language Models

Kabir Man Singh | 24059480

Neural Networks and Deep Learning

Introduction & Problem

- Combining computer vision and NLP to generate meaningful descriptions of images
- Why it matters: Assistive technology, content retrieval, social media automation
- Challenge: Bridging the semantic gap between visual features and natural language
- Goal: Build an end-to-end automatic image caption generation system using deep learning

Dataset & EDA

- Database:Flickr8k
- 8,091 Images | 40,455 Captions | 5 Captions per Image
- Average caption length: 12 words (range 8–15 words)
- Image dimensions: ~400×450px average
- Vocabulary: 5,000 tokens

Dataset & EDA

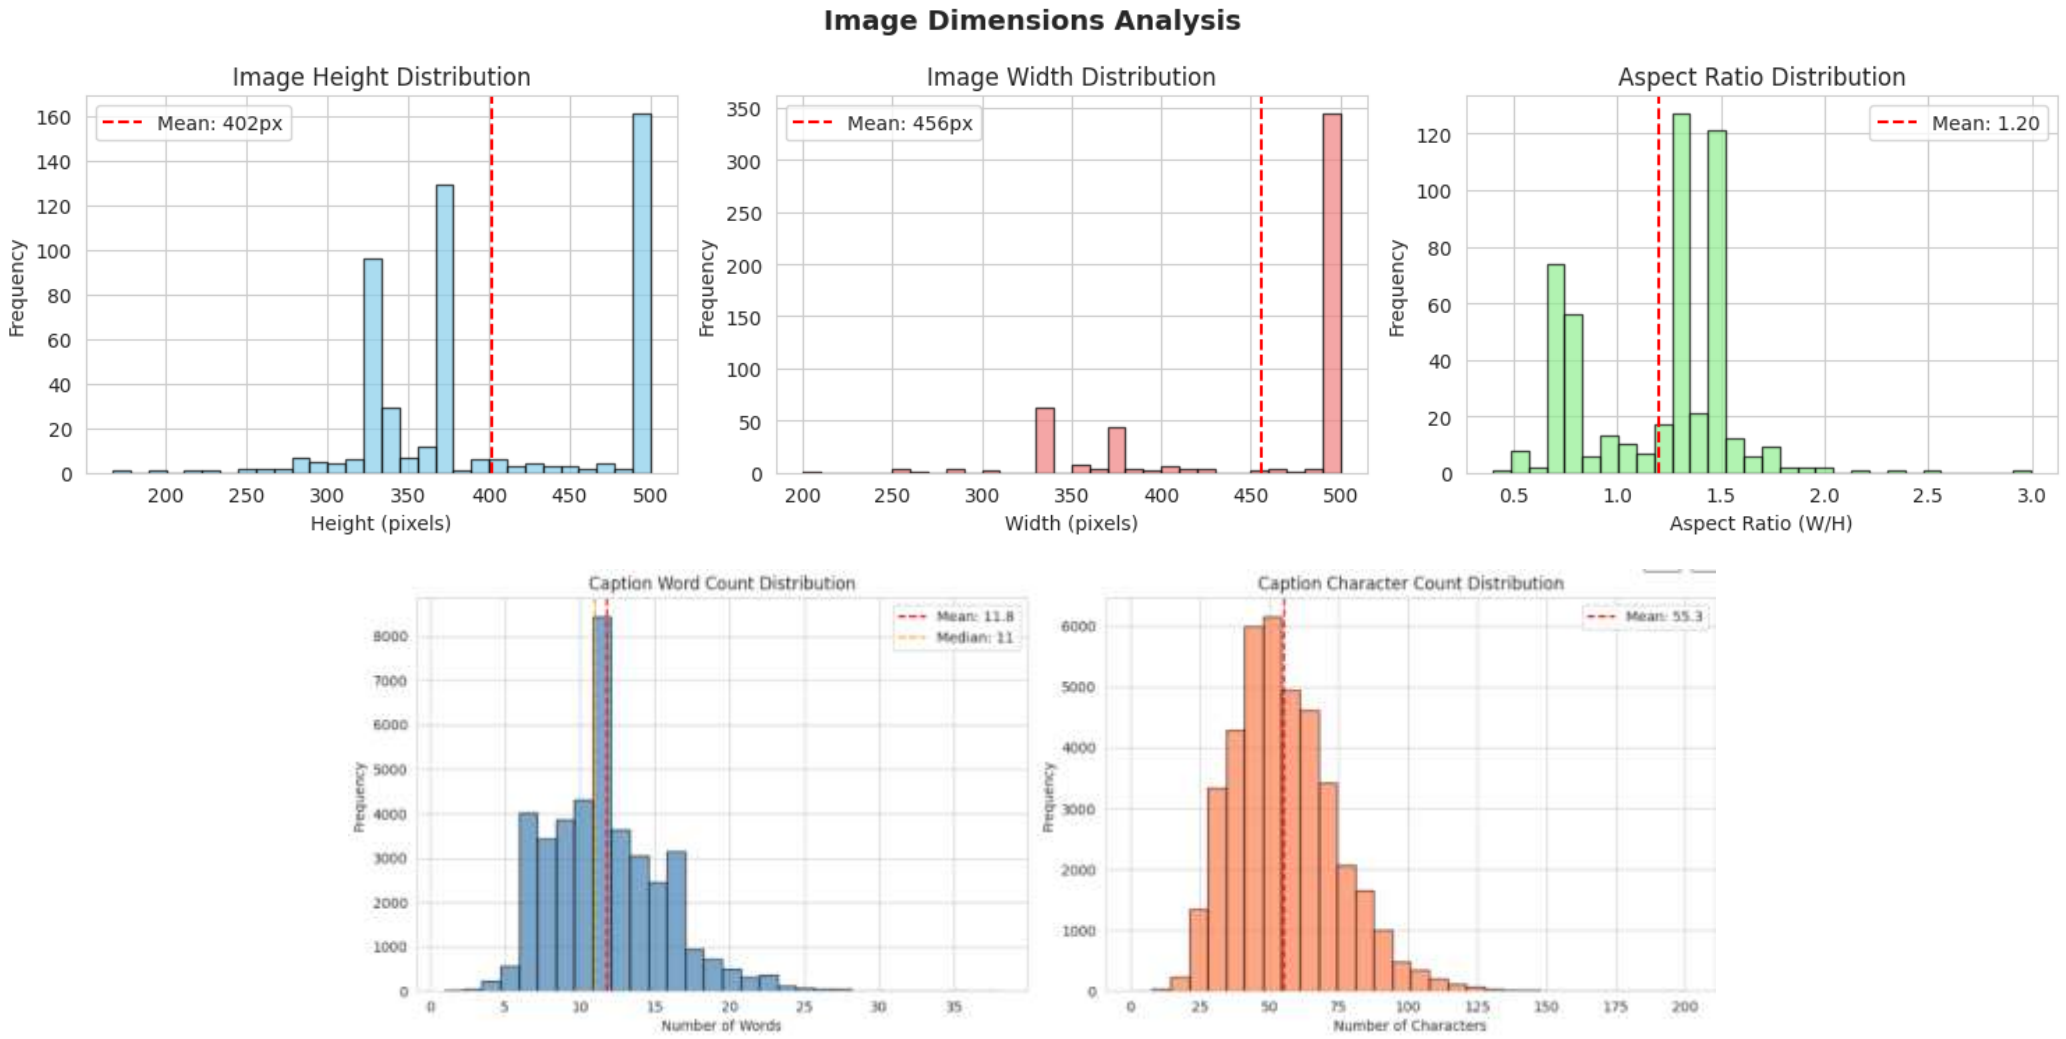
Image Dimensions Analysis
Image Height Distribution
Mean: 402px
Frequency
Height (pixels)
Image Width Distribution
Mean: 456px
Frequency
Width (pixels)
Aspect Ratio Distribution
Mean: 1.20
Frequency
Aspect Ratio (W/H)
Caption Word Count Distribution
Mean: 11.8
Median: 11
Frequency
Number of Words
Caption Character Count Distribution
Mean: 55.3
Frequency
Number of Characters

Methodology & Architecture

- CNN Encoder: EfficientNetB0 (frozen, pre-trained on ImageNet) → 7×7×1280 feature map → reshaped to 49×1280 → projected to 256D
- Transformer Encoder: 1 layer, 4-head multi-head self-attention + FFN
- Transformer Decoder: Masked self-attention + cross-attention → Dense output layer (5,000 vocab)

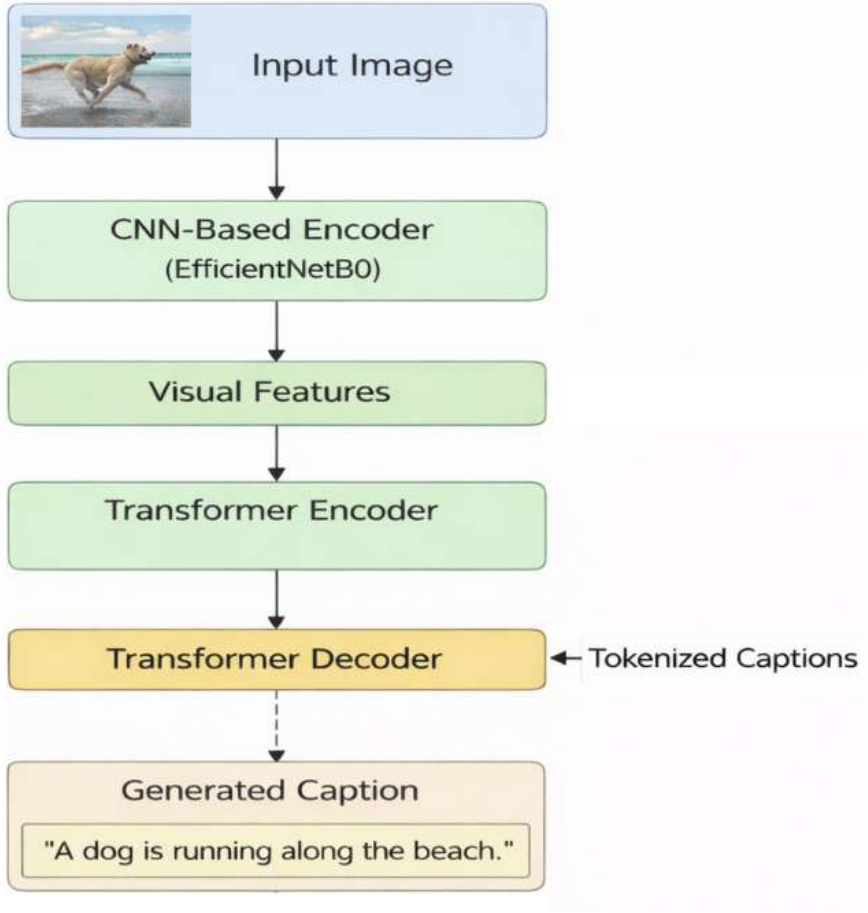

Implementation & Configuration

- Optimizer: Adam (learning rate 0.001)
- Number of Layers: Transformer Encoder (1 layer), Transformer Decoder (1 layer)
- Total Trainable Parameters: 10,495,275
- Total Epochs Trained: 10
- Dataset: Flickr8k (8,091 images, 40,455 captions)
- Loss Function: Sparse Categorical Crossentropy
- Activation Functions: ReLU (FFN), Softmax (output)

Training Results

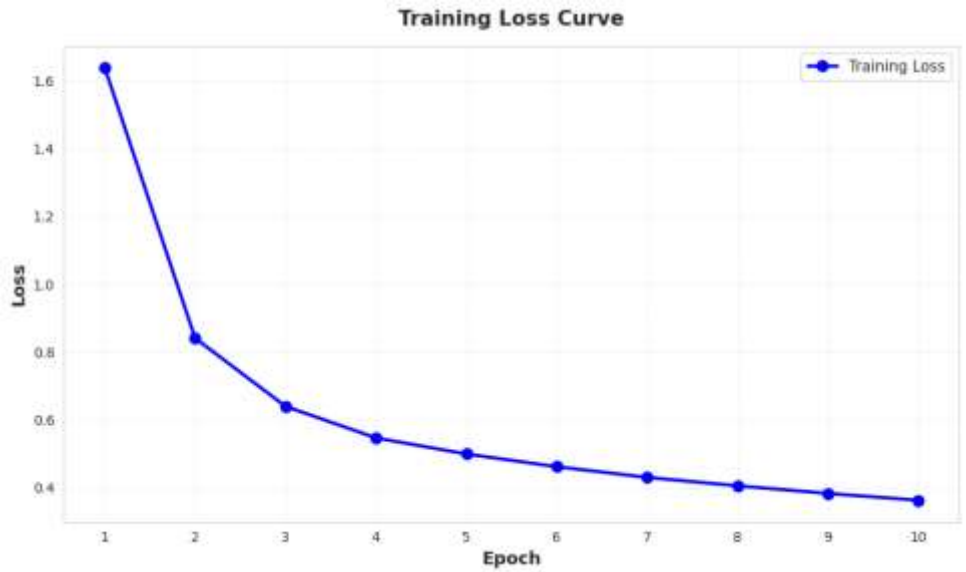

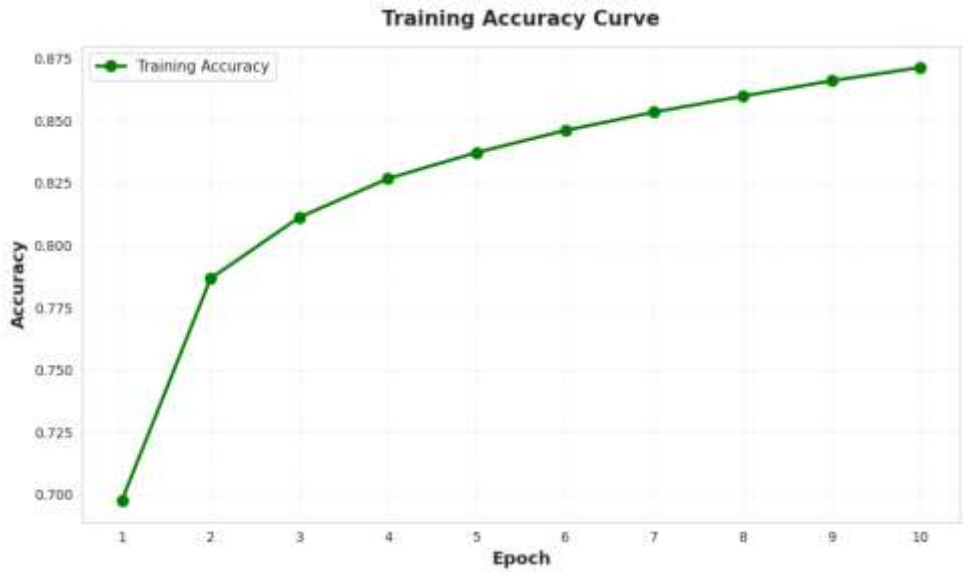

Metric	Initial	Final	Best	Best Epoch	Improvement
Loss	1.6388	0.3632	0.3632	10	1.2756
Accuracy	0.6977	0.8714	0.8714	10	0.1736

Visual Evidence

```
Loading model configuration and weights...
Model weights loaded successfully.

Generating caption for: car.jpg

PREDICTED CAPTION:
→ a blue car parked next to a road
(venv) PS C:\Project\Image-Captioning_local>
```

```
Generating caption for: dog.jpg

PREDICTED CAPTION:
→ a dog laying on the grass with a frisbee in its mouth
(venv) PS C:\Project\Image-Captioning_local>
```

Strengths & Weaknesses

Strengths:

- Consistent convergence without overfitting
- Effective transfer learning with EfficientNetB0
- Successful visual-linguistic feature integration
- Resource-efficient (CPU-only training)

Weakness:

- No validation split for generalization testing
- Restricted vocabulary (5,000 tokens)
- No BLEU score evaluation
- Limited performance on complex scenes

Conclusion

- Successfully built end-to-end image captioning system
- Successfully generated captions using deep learning
- Demonstrated practical applications of CV and NLP
- Achieved 87.14% accuracy with 77.8% loss reduction

Future Work

- Use larger datasets (MS COCO)
- Implement train-validation-test splits
- Attention visualization for interpretability

Thank You!